

RECENT FORM

Mean or median? It depends on the job.

For four-gameweek ranks, they are tied.

For exact forecasts, rare hauls change the answer.

Points are lumpy

STEP 1


Many players score zero or one in a gameweek. A few scores are very large.

Even among full-match players, the mean sits above the median.

That gap is made by rare, **one-off big hauls**.

Rare big hauls pull the mean away from the median.

A single gameweek score is right-skewed: mostly small, occasionally huge



The difference appears even among players who played.

Four-game summaries disagree


STEP 2

Over four gameweeks, median stays spiky at low whole numbers.

Mean is smoother and shifted higher.

One big haul can pull up the average without changing the median at all.

Mean and median tell different four-game stories.



The actual recent-form feature for likely starters.

For ranking, they are close

STEP 3


Both past-four-game signals have a wide relationship with the next four games.

Their trend lines are **barely distinguishable** at this horizon.

The ranking figures are 0.203 for mean and 0.215 for median.

At four gameweeks, mean and median rank similarly.

Past form vs. what actually happened next: a real but noisy link either way



A real but weak recent-form signal either way.

The winner flips with the task

STEP 4


For ranking, mean and median tie through six games.
Mean pulls ahead at seven and eight.

For exact-point forecasts, median is more accurate at four and six games.

There is no single best summary without naming the job first.

Ranking and exact forecasts reward different choices.

Ranking and forecast-error tell different stories across window length



Results change with the horizon and outcome.

Midfielders drive the gap

STEP 5



Pooling positions hides the main pattern.

Midfielders create almost all of the median advantage for forecast error.

Defenders, forwards, and keepers show little clear pattern.

The median edge is mostly a midfielder story.

The pooled numbers are a compromise — midfielders drive most of the median story



Position mix changes the size of the effect.

The problem is the haul


STEP 6

Cap past scores at 10 points and almost all of median's forecast-error edge disappears.

That shows mean is not wrong. It is just exposed to rare hauls that do not repeat.

A capped mean can catch back up, or win.

Capping big hauls **erases almost all** of the median edge.



The mechanism is unguarded one-off scores.

Choose the stat for the job

WHAT TO DO

- For four-gameweek transfer ranks, **keep mean-based ranking**.
- Do not switch rules for a statistical tie.
- For exact forecasts, cap extreme past scores before averaging.
- Treat a capped mean as a next step, not a tested final answer.

A context, not a verdict

CAVEATS + METHOD

This is one 2025/26 season. Other seasons are not yet tested.

These results are links, not causes of future points.

The likely-starter gate and position mix affect the size of the result.

We compared walk-forward four-to-eight-game windows with gameweek-block resampling.

At N=4, keep the mean for ranks. Guard against hauls for exact forecasts.